

Assignment: Artificial Neural Networks & Deep Learning

Daniel Fuentes, Marc Valenti

5/10/2018

Contents

NN architectures	1
CNN architectures:	49

Data sets: Caltech 101 Description

In the web page: http://www.vision.caltech.edu/Image_Datasets/Caltech101/> there are pictures of objects belonging to 101 categories. About 40 to 800 images per category. Most categories have about 50 images. Collected in September 2003 by Fei-Fei Li, Marco Andreetto, and Marc 'Aurelio Ranzato. The size of each image is roughly 300 x 200 pixels.

NN architectures

- Select five categories from caltech101 collection. Try to select balanced categories with representative number of images.

```
#only replace the path_images string with where the dataset of pictures from http://www.vision.caltech.edu/Image_Datasets/Caltech101/
#path_images <- "C:/Users/DanEscario/Desktop/MESIO/Statistical learning/Pr?ctiques/XarxaNeural/101_ObjectCategories/"
path_images <- "~/Desktop/101_ObjectCategories/"

categories <- list.files(path=path_images)

number_images <- rep(NA,length(categories))

for (i in 1:length(categories)){
  number_images[i]<-list.files(path=paste0(path_images,categories[i])) %>%
  length()
}

selected_categories <- categories[number_images%in%(86:88)]
```

After inspecting the number of images of every category, we have ended up choosing *helicopter*, *kangaroo*, *menorah*, *starfish*, *trilobite*, which have 88, 86, 87, 86, 86 images each.

```
# input_size <- 128 # Width and height of images supplied to the input of a neural network
# epochs <- 30 # The number of epochs
# batch_size <- 16 # Batch size
# orig_width <- 300 # Original image width
# orig_height <- 200 # Height of source images
# train_samples <- 5088 # Size of training sample
# train_index <- sample (1: train_samples, round (train_samples * 0.8)) # 80%
# val_index <- c (1: train_samples) [- train_index]
```

```

## Folders with pictures
images_dir <- path_images
# masks_dir <- "input / train_masks /"

image_to_grey <- function(image_file,
                           target_width = 32,
                           target_height = 32) {
  img <- image_read(image_file) %>%
    image_convert(colorspace = "gray") %>%
    image_scale(paste0(target_width, "x", target_height, "!"))
  img
}

img2arr <- function(image) {
  result <- as.numeric(image[[1]])
  return(result)
}

#llegeix i escala la imatge
img <- image_to_grey(paste0(path_images, "bonsai/image_0013.jpg"))
#converteix la imatge en un array
img_test <- as.numeric(image_to_grey(paste0(path_images, "bonsai/image_0012.jpg"))[[1]])

```

Implements a classifier fed with the HOG descriptors that learns the category to which a given image belongs. In HOG use 3 cells and 9 orientations. Compare the performance of two NN architectures: 1. Layer 1: 81 units (input layer), Layer 2: 10 units (hidden layer), Layer 3: 5 units (output layer), 2. Layer 1: 81 units (input layer), Layer 2: 50 units (hidden layer), Layer 3: 25 units (hidden layer), Layer 4: 5 units (output layer).

```

obtain_images_to_HOG <- function(x,
                                vector_categories,
                                number_cells=3,
                                number_orientations=9
                                ){

  #gets the list of the files in a given directory first.
  imgs <- list.files(paste0(images_dir, vector_categories[x]),
                    pattern = ".jpg",
                    full.names = TRUE)

  #creates an auxiliary matrix where we will place the results.
  #every row is a image. it first contains the label, and then the 1024 vector
  dt_img <- as.data.frame(rep(vector_categories[x], number_images[categories==vector_categories[x]]))
  colnames(dt_img) <- "label"
  dt_img[, 2:82] <- 0

  #we iterate on all the images
  for(i in 1:length(imgs)) {
    if(length(dim(readImage(imgs[i])))==3){
      img <- readImage(imgs[i]) %>%
        rgb2gray() %>%
        resizeImage(width = 32,

```

```

        height = 32,
        method = "nearest") %>%
    HOG(cells=number_cells,
        orientations=number_orientations)
dt_img[i,2:82] <- img
}
}

return(dt_img)
}

dt_HOG <- rbind(obtain_images_to_HOG(1,vector_categories=selected_categories),
    obtain_images_to_HOG(2,selected_categories),
    obtain_images_to_HOG(3,selected_categories),
    obtain_images_to_HOG(4,selected_categories),
    obtain_images_to_HOG(5,selected_categories))

trainIndex <- createDataPartition(dt_HOG$label, p=.7, list=F)
dmy <- dummyVars(" ~ label", data = dt_HOG)
dmy_aux <- data.frame(predict(dmy, newdata = dt_HOG))
y_dt_HOG<-dt_HOG$label
dmy_aux <- data.frame(apply(dmy_aux, 2, function(x) (as.numeric(as.character(x)))))
dt_HOG <- cbind(dt_HOG,dmy_aux)
dt_HOG$label<-NULL

resp_var_HOG <- names(select(dmy_aux,starts_with("label")))
xnam <- names(dt_HOG[!names(dt_HOG)%in%resp_var_HOG])

(fmla <- as.formula(paste(paste(resp_var_HOG,collapse="+"), "~", paste(xnam, collapse= "+"))))

## label.helicopter + label.kangaroo + label.menorah + label.starfish +
##     label.trilobite ~ V2 + V3 + V4 + V5 + V6 + V7 + V8 + V9 +
##     V10 + V11 + V12 + V13 + V14 + V15 + V16 + V17 + V18 + V19 +
##     V20 + V21 + V22 + V23 + V24 + V25 + V26 + V27 + V28 + V29 +
##     V30 + V31 + V32 + V33 + V34 + V35 + V36 + V37 + V38 + V39 +
##     V40 + V41 + V42 + V43 + V44 + V45 + V46 + V47 + V48 + V49 +
##     V50 + V51 + V52 + V53 + V54 + V55 + V56 + V57 + V58 + V59 +
##     V60 + V61 + V62 + V63 + V64 + V65 + V66 + V67 + V68 + V69 +
##     V70 + V71 + V72 + V73 + V74 + V75 + V76 + V77 + V78 + V79 +
##     V80 + V81 + V82

dt_HOG_train <- dt_HOG[trainIndex, ]
dt_HOG_test <- dt_HOG[-trainIndex, ]

HOG_grid_10 = expand.grid(layer1 = 10,
    layer2 = 0,
    layer3 = 0)

HOG_10_CPU_time <- Sys.time()

HOG.fit_10 = neuralnet(fmla,
    data = dt_HOG_train,
    hidden= 10,
    linear.output=FALSE)

```

```

HOG_10_CPU_time <- Sys.time() - HOG_10_CPU_time

HOG_grid_50 = expand.grid(layer1 = 50,
                          layer2 = 25,
                          layer3=5)

HOG_50_CPU_time <- Sys.time()

HOG_fit_50 = neuralnet(fmla,
                      data = dt_HOG_train,
                      hidden= c(50, 25),
                      linear.output=FALSE)

HOG_50_CPU_time <- Sys.time() - HOG_50_CPU_time

maxidx <- function(row) {
  return(which(row == max(row)))
}

# Plot model and predict
#plotnet(HOG_fit_10) #dona error: "Error in plotnet(HOG_fit): object HOG_fit not found"
train_pred_10 <- compute(HOG_fit_10, dt_HOG_train[,!names(dt_HOG_train) %in% resp_var_HOG])$net.result %>%
  apply(1, maxidx)

test_pred_10 <- compute(HOG_fit_10, dt_HOG_test[,!names(dt_HOG_test) %in% resp_var_HOG])$net.result %>%
  apply(1, maxidx)

#plotnet(HOG_fit_50)
train_pred_50 <- compute(HOG_fit_50, dt_HOG_train[,!names(dt_HOG_train) %in% resp_var_HOG])$net.result %>%
  apply(1, maxidx)
test_pred_50 <- compute(HOG_fit_50, dt_HOG_test[,!names(dt_HOG_test) %in% resp_var_HOG])$net.result %>%
  apply(1, maxidx)

```

3. Reshape grayscale images to 32x32. Vectorize images by rows, to get a 1024-vector representation of each image. Implement a NN with two hidden layers that uses such a vectors as input data for learning image categories. Tune the number of neurons in each hidden layer using a grid search procedure.

```

#let's create the dataset first:
#the function obtain_images, would allow us to do this

obtain_images <- function(x,vector_categories){

  #gets the list of the files in a given directory first.
  imgs <- list.files(paste0(images_dir,selected_categories[x]),
                    pattern = ".jpg",
                    full.names = TRUE)

  #creates an auxiliary matrix where we will place the results.
  #every row is a image. it first contains the label, and then the 1024 vector
  dt_img <- as.data.frame(rep(vector_categories[x],number_images[categories==vector_categories[x]]))
  colnames(dt_img)<-"label"
  dt_img[,2:1025]<-0

  #we iterate on all the images

```

```


for(i in 1:length(imgs)) {
  img <- image_to_grey(imgs[i])
  dt_img[i,2:1025] <- as.vector(img2arr(img))
}

return(dt_img)
}

#we create the dataset by row binding the different categories
dt_ex3 <- rbind(obtain_images(1,selected_categories),
               obtain_images(2,selected_categories),
               obtain_images(3,selected_categories),
               obtain_images(4,selected_categories),
               obtain_images(5,selected_categories))

#let's see if that works!
matrix(unlist(dt_ex3[30,-1]), nrow = 32, byrow = TRUE) %>%
  image(col=grey.colors(255))

```



```

image_read(paste0(list.files(paste0(images_dir,selected_categories[1]),
                             pattern = ".jpg",
                             full.names = TRUE))[30])

```



#It does! it just acts as a mirror.

#write.csv2(dt,file="data_free_categories_nn.csv")

```
trainIndex <- createDataPartition(dt_ex3$label, p=.7, list=F)
dt_ex3_train <- dt_ex3[trainIndex, ]
dt_ex3_test <- dt_ex3[-trainIndex, ]

xnam <- names(dt_ex3_train[-1])
resp_var <- paste0(levels(dt_ex3_train$label),collapse=" + ")
(fmla <- as.formula(paste(resp_var, "~", paste(xnam, collapse= "+"))))
```

```
## helicopter + kangaroo + menorah + starfish + trilobite ~ V2 +
##      V3 + V4 + V5 + V6 + V7 + V8 + V9 + V10 + V11 + V12 + V13 +
##      V14 + V15 + V16 + V17 + V18 + V19 + V20 + V21 + V22 + V23 +
##      V24 + V25 + V26 + V27 + V28 + V29 + V30 + V31 + V32 + V33 +
##      V34 + V35 + V36 + V37 + V38 + V39 + V40 + V41 + V42 + V43 +
##      V44 + V45 + V46 + V47 + V48 + V49 + V50 + V51 + V52 + V53 +
##      V54 + V55 + V56 + V57 + V58 + V59 + V60 + V61 + V62 + V63 +
##      V64 + V65 + V66 + V67 + V68 + V69 + V70 + V71 + V72 + V73 +
##      V74 + V75 + V76 + V77 + V78 + V79 + V80 + V81 + V82 + V83 +
##      V84 + V85 + V86 + V87 + V88 + V89 + V90 + V91 + V92 + V93 +
##      V94 + V95 + V96 + V97 + V98 + V99 + V100 + V101 + V102 +
##      V103 + V104 + V105 + V106 + V107 + V108 + V109 + V110 + V111 +
##      V112 + V113 + V114 + V115 + V116 + V117 + V118 + V119 + V120 +
##      V121 + V122 + V123 + V124 + V125 + V126 + V127 + V128 + V129 +
##      V130 + V131 + V132 + V133 + V134 + V135 + V136 + V137 + V138 +
##      V139 + V140 + V141 + V142 + V143 + V144 + V145 + V146 + V147 +
##      V148 + V149 + V150 + V151 + V152 + V153 + V154 + V155 + V156 +
##      V157 + V158 + V159 + V160 + V161 + V162 + V163 + V164 + V165 +
##      V166 + V167 + V168 + V169 + V170 + V171 + V172 + V173 + V174 +
##      V175 + V176 + V177 + V178 + V179 + V180 + V181 + V182 + V183 +
##      V184 + V185 + V186 + V187 + V188 + V189 + V190 + V191 + V192 +
##      V193 + V194 + V195 + V196 + V197 + V198 + V199 + V200 + V201 +
##      V202 + V203 + V204 + V205 + V206 + V207 + V208 + V209 + V210 +
```

```

## V211 + V212 + V213 + V214 + V215 + V216 + V217 + V218 + V219 +
## V220 + V221 + V222 + V223 + V224 + V225 + V226 + V227 + V228 +
## V229 + V230 + V231 + V232 + V233 + V234 + V235 + V236 + V237 +
## V238 + V239 + V240 + V241 + V242 + V243 + V244 + V245 + V246 +
## V247 + V248 + V249 + V250 + V251 + V252 + V253 + V254 + V255 +
## V256 + V257 + V258 + V259 + V260 + V261 + V262 + V263 + V264 +
## V265 + V266 + V267 + V268 + V269 + V270 + V271 + V272 + V273 +
## V274 + V275 + V276 + V277 + V278 + V279 + V280 + V281 + V282 +
## V283 + V284 + V285 + V286 + V287 + V288 + V289 + V290 + V291 +
## V292 + V293 + V294 + V295 + V296 + V297 + V298 + V299 + V300 +
## V301 + V302 + V303 + V304 + V305 + V306 + V307 + V308 + V309 +
## V310 + V311 + V312 + V313 + V314 + V315 + V316 + V317 + V318 +
## V319 + V320 + V321 + V322 + V323 + V324 + V325 + V326 + V327 +
## V328 + V329 + V330 + V331 + V332 + V333 + V334 + V335 + V336 +
## V337 + V338 + V339 + V340 + V341 + V342 + V343 + V344 + V345 +
## V346 + V347 + V348 + V349 + V350 + V351 + V352 + V353 + V354 +
## V355 + V356 + V357 + V358 + V359 + V360 + V361 + V362 + V363 +
## V364 + V365 + V366 + V367 + V368 + V369 + V370 + V371 + V372 +
## V373 + V374 + V375 + V376 + V377 + V378 + V379 + V380 + V381 +
## V382 + V383 + V384 + V385 + V386 + V387 + V388 + V389 + V390 +
## V391 + V392 + V393 + V394 + V395 + V396 + V397 + V398 + V399 +
## V400 + V401 + V402 + V403 + V404 + V405 + V406 + V407 + V408 +
## V409 + V410 + V411 + V412 + V413 + V414 + V415 + V416 + V417 +
## V418 + V419 + V420 + V421 + V422 + V423 + V424 + V425 + V426 +
## V427 + V428 + V429 + V430 + V431 + V432 + V433 + V434 + V435 +
## V436 + V437 + V438 + V439 + V440 + V441 + V442 + V443 + V444 +
## V445 + V446 + V447 + V448 + V449 + V450 + V451 + V452 + V453 +
## V454 + V455 + V456 + V457 + V458 + V459 + V460 + V461 + V462 +
## V463 + V464 + V465 + V466 + V467 + V468 + V469 + V470 + V471 +
## V472 + V473 + V474 + V475 + V476 + V477 + V478 + V479 + V480 +
## V481 + V482 + V483 + V484 + V485 + V486 + V487 + V488 + V489 +
## V490 + V491 + V492 + V493 + V494 + V495 + V496 + V497 + V498 +
## V499 + V500 + V501 + V502 + V503 + V504 + V505 + V506 + V507 +
## V508 + V509 + V510 + V511 + V512 + V513 + V514 + V515 + V516 +
## V517 + V518 + V519 + V520 + V521 + V522 + V523 + V524 + V525 +
## V526 + V527 + V528 + V529 + V530 + V531 + V532 + V533 + V534 +
## V535 + V536 + V537 + V538 + V539 + V540 + V541 + V542 + V543 +
## V544 + V545 + V546 + V547 + V548 + V549 + V550 + V551 + V552 +
## V553 + V554 + V555 + V556 + V557 + V558 + V559 + V560 + V561 +
## V562 + V563 + V564 + V565 + V566 + V567 + V568 + V569 + V570 +
## V571 + V572 + V573 + V574 + V575 + V576 + V577 + V578 + V579 +
## V580 + V581 + V582 + V583 + V584 + V585 + V586 + V587 + V588 +
## V589 + V590 + V591 + V592 + V593 + V594 + V595 + V596 + V597 +
## V598 + V599 + V600 + V601 + V602 + V603 + V604 + V605 + V606 +
## V607 + V608 + V609 + V610 + V611 + V612 + V613 + V614 + V615 +
## V616 + V617 + V618 + V619 + V620 + V621 + V622 + V623 + V624 +
## V625 + V626 + V627 + V628 + V629 + V630 + V631 + V632 + V633 +
## V634 + V635 + V636 + V637 + V638 + V639 + V640 + V641 + V642 +
## V643 + V644 + V645 + V646 + V647 + V648 + V649 + V650 + V651 +
## V652 + V653 + V654 + V655 + V656 + V657 + V658 + V659 + V660 +
## V661 + V662 + V663 + V664 + V665 + V666 + V667 + V668 + V669 +
## V670 + V671 + V672 + V673 + V674 + V675 + V676 + V677 + V678 +
## V679 + V680 + V681 + V682 + V683 + V684 + V685 + V686 + V687 +
## V688 + V689 + V690 + V691 + V692 + V693 + V694 + V695 + V696 +

```

```

##      V697 + V698 + V699 + V700 + V701 + V702 + V703 + V704 + V705 +
##      V706 + V707 + V708 + V709 + V710 + V711 + V712 + V713 + V714 +
##      V715 + V716 + V717 + V718 + V719 + V720 + V721 + V722 + V723 +
##      V724 + V725 + V726 + V727 + V728 + V729 + V730 + V731 + V732 +
##      V733 + V734 + V735 + V736 + V737 + V738 + V739 + V740 + V741 +
##      V742 + V743 + V744 + V745 + V746 + V747 + V748 + V749 + V750 +
##      V751 + V752 + V753 + V754 + V755 + V756 + V757 + V758 + V759 +
##      V760 + V761 + V762 + V763 + V764 + V765 + V766 + V767 + V768 +
##      V769 + V770 + V771 + V772 + V773 + V774 + V775 + V776 + V777 +
##      V778 + V779 + V780 + V781 + V782 + V783 + V784 + V785 + V786 +
##      V787 + V788 + V789 + V790 + V791 + V792 + V793 + V794 + V795 +
##      V796 + V797 + V798 + V799 + V800 + V801 + V802 + V803 + V804 +
##      V805 + V806 + V807 + V808 + V809 + V810 + V811 + V812 + V813 +
##      V814 + V815 + V816 + V817 + V818 + V819 + V820 + V821 + V822 +
##      V823 + V824 + V825 + V826 + V827 + V828 + V829 + V830 + V831 +
##      V832 + V833 + V834 + V835 + V836 + V837 + V838 + V839 + V840 +
##      V841 + V842 + V843 + V844 + V845 + V846 + V847 + V848 + V849 +
##      V850 + V851 + V852 + V853 + V854 + V855 + V856 + V857 + V858 +
##      V859 + V860 + V861 + V862 + V863 + V864 + V865 + V866 + V867 +
##      V868 + V869 + V870 + V871 + V872 + V873 + V874 + V875 + V876 +
##      V877 + V878 + V879 + V880 + V881 + V882 + V883 + V884 + V885 +
##      V886 + V887 + V888 + V889 + V890 + V891 + V892 + V893 + V894 +
##      V895 + V896 + V897 + V898 + V899 + V900 + V901 + V902 + V903 +
##      V904 + V905 + V906 + V907 + V908 + V909 + V910 + V911 + V912 +
##      V913 + V914 + V915 + V916 + V917 + V918 + V919 + V920 + V921 +
##      V922 + V923 + V924 + V925 + V926 + V927 + V928 + V929 + V930 +
##      V931 + V932 + V933 + V934 + V935 + V936 + V937 + V938 + V939 +
##      V940 + V941 + V942 + V943 + V944 + V945 + V946 + V947 + V948 +
##      V949 + V950 + V951 + V952 + V953 + V954 + V955 + V956 + V957 +
##      V958 + V959 + V960 + V961 + V962 + V963 + V964 + V965 + V966 +
##      V967 + V968 + V969 + V970 + V971 + V972 + V973 + V974 + V975 +
##      V976 + V977 + V978 + V979 + V980 + V981 + V982 + V983 + V984 +
##      V985 + V986 + V987 + V988 + V989 + V990 + V991 + V992 + V993 +
##      V994 + V995 + V996 + V997 + V998 + V999 + V1000 + V1001 +
##      V1002 + V1003 + V1004 + V1005 + V1006 + V1007 + V1008 + V1009 +
##      V1010 + V1011 + V1012 + V1013 + V1014 + V1015 + V1016 + V1017 +
##      V1018 + V1019 + V1020 + V1021 + V1022 + V1023 + V1024 + V1025

```

```

my.grid <- expand.grid(size=seq(from=1,to=10,by=2),
                      decay=seq(from=0.1,to=1,by=0.2))
ex3_grid_search_time<-Sys.time()
ex3.fit <- train(label ~ .,
                data = dt_ex3_train,
                method = "nnet",
                tuneGrid = my.grid,
                verbose=FALSE,
                preProc = c('center', 'scale', 'knnImpute', 'pca'),
                trControl = trainControl(method = "cv",
                                         verboseIter = FALSE,
                                         number = 5,
                                         returnData = FALSE)

                # ,MaxNWts=1036
                )

```

```
## # weights: 85
```




```

## iter 50 value 173.545654
## iter 60 value 162.719399
## iter 70 value 155.870170
## iter 80 value 148.626149
## iter 90 value 145.409689
## iter 100 value 144.297627
## final value 144.297627
## stopped after 100 iterations
## # weights: 716
## initial value 595.572725
## iter 10 value 331.362170
## iter 20 value 250.856571
## iter 30 value 195.719995
## iter 40 value 163.772559
## iter 50 value 151.357016
## iter 60 value 143.996295
## iter 70 value 138.247261
## iter 80 value 134.673976
## iter 90 value 133.142627
## iter 100 value 132.064214
## final value 132.064214
## stopped after 100 iterations
## # weights: 621
## initial value 647.664528
## iter 10 value 368.687520
## iter 20 value 304.269749
## iter 30 value 262.423579
## iter 40 value 217.116086
## iter 50 value 184.993070
## iter 60 value 164.846343
## iter 70 value 154.594602
## iter 80 value 143.777923
## iter 90 value 139.169516
## iter 100 value 136.392598
## final value 136.392598
## stopped after 100 iterations
ex3_grid_search_time<-Sys.time() - ex3_grid_search_time

plot(ex3.fit)

```



Doing the grid search procedure with a 5 fold CV has taken 1.5308816512. The CPU time in this case will be a function of how big our grid.search is and how many folds the CV has. Therefore, we now estimate the new CPU time with the best tune: 7 as size and 0.5 as decay.

```
my.grid <- expand.grid(size=ex3.fit$bestTune$size,
                      decay=ex3.fit$bestTune$decay)

ex3_fit_time<-Sys.time()
ex3.fit <- train(label ~ .,
                 data = dt_ex3_train,
                 method = "nnet",
                 tuneGrid = my.grid,
                 verbose=FALSE,
                 preProc = c('center', 'scale', 'knnImpute', 'pca'),
                 trControl = trainControl(method = "cv",
                                           verboseIter = FALSE,
                                           number = 10,
                                           returnData = FALSE)

# ,MaxNWts=1036
)
```

```
## # weights: 607
## initial value 559.015309
## iter 10 value 330.779221
## iter 20 value 281.536276
## iter 30 value 229.926281
## iter 40 value 182.780938
```

```

## iter 50 value 155.450435
## iter 60 value 145.249051
## iter 70 value 131.244724
## iter 80 value 124.876203
## iter 90 value 121.602509
## iter 100 value 120.441214
## final value 120.441214
## stopped after 100 iterations
## # weights: 593
## initial value 561.952372
## iter 10 value 318.815677
## iter 20 value 262.843930
## iter 30 value 224.132340
## iter 40 value 194.248426
## iter 50 value 177.037208
## iter 60 value 165.775516
## iter 70 value 155.380407
## iter 80 value 150.097929
## iter 90 value 146.006412
## iter 100 value 143.853871
## final value 143.853871
## stopped after 100 iterations
## # weights: 579
## initial value 619.190388
## iter 10 value 353.483766
## iter 20 value 297.533552
## iter 30 value 252.116005
## iter 40 value 213.262950
## iter 50 value 197.084655
## iter 60 value 182.811745
## iter 70 value 164.731693
## iter 80 value 152.195225
## iter 90 value 143.907690
## iter 100 value 138.547114
## final value 138.547114
## stopped after 100 iterations
## # weights: 586
## initial value 523.024303
## iter 10 value 323.177873
## iter 20 value 266.970861
## iter 30 value 216.298884
## iter 40 value 181.066403
## iter 50 value 156.439151
## iter 60 value 142.749586
## iter 70 value 134.291925
## iter 80 value 127.016980
## iter 90 value 123.074999
## iter 100 value 120.944222
## final value 120.944222
## stopped after 100 iterations
## # weights: 600
## initial value 576.970862
## iter 10 value 389.825347
## iter 20 value 319.780232

```

```

## iter 30 value 268.733501
## iter 40 value 220.299625
## iter 50 value 188.195881
## iter 60 value 160.117232
## iter 70 value 146.437586
## iter 80 value 135.667292
## iter 90 value 131.264236
## iter 100 value 129.446844
## final value 129.446844
## stopped after 100 iterations
## # weights: 600
## initial value 549.236699
## iter 10 value 321.680284
## iter 20 value 255.093438
## iter 30 value 211.477523
## iter 40 value 174.774963
## iter 50 value 150.439103
## iter 60 value 141.708056
## iter 70 value 130.085627
## iter 80 value 125.884751
## iter 90 value 123.516667
## iter 100 value 121.999893
## final value 121.999893
## stopped after 100 iterations
## # weights: 586
## initial value 557.140801
## iter 10 value 342.120900
## iter 20 value 284.489204
## iter 30 value 242.641977
## iter 40 value 209.161666
## iter 50 value 183.686934
## iter 60 value 164.344607
## iter 70 value 152.345142
## iter 80 value 142.383273
## iter 90 value 134.208212
## iter 100 value 131.313107
## final value 131.313107
## stopped after 100 iterations
## # weights: 593
## initial value 501.967916
## iter 10 value 315.481536
## iter 20 value 271.561921
## iter 30 value 220.536532
## iter 40 value 191.682009
## iter 50 value 166.259417
## iter 60 value 152.692319
## iter 70 value 139.912162
## iter 80 value 130.452532
## iter 90 value 126.773669
## iter 100 value 125.174035
## final value 125.174035
## stopped after 100 iterations
## # weights: 600
## initial value 515.652097

```

```
## iter 10 value 331.188144
## iter 20 value 287.206404
## iter 30 value 233.344038
## iter 40 value 193.987888
## iter 50 value 171.576418
## iter 60 value 154.184445
## iter 70 value 140.245486
## iter 80 value 127.058210
## iter 90 value 120.888641
## iter 100 value 118.134358
## final value 118.134358
## stopped after 100 iterations
## # weights: 593
## initial value 538.039721
## iter 10 value 338.931176
## iter 20 value 286.719812
## iter 30 value 235.833928
## iter 40 value 191.486994
## iter 50 value 168.345129
## iter 60 value 156.983503
## iter 70 value 144.102982
## iter 80 value 136.325986
## iter 90 value 133.412310
## iter 100 value 130.447274
## final value 130.447274
## stopped after 100 iterations
## # weights: 621
## initial value 580.892030
## iter 10 value 355.620986
## iter 20 value 304.526695
## iter 30 value 255.737697
## iter 40 value 203.464396
## iter 50 value 183.018423
## iter 60 value 167.813138
## iter 70 value 158.789683
## iter 80 value 151.289120
## iter 90 value 148.549402
## iter 100 value 145.918786
## final value 145.918786
## stopped after 100 iterations
```

```
ex3_fit_time<-Sys.time() - ex3_fit_time
```

```
train_pred_ex3 <- predict(ex3.fit, newdata=dt_ex3_train, type="raw")
test_pred_ex3 <- predict(ex3.fit, newdata=dt_ex3_test, type="raw")
```

4. Compare these three NN model at least in two characteristics: the performance and CPU time.

```
taulahog10<-(table(test_pred_10, max.col(dt_HOG_test[,names(dt_HOG_test)%in%resp_var_HOG])));
cmhog10<-confusionMatrix(taulahog10);cmhog10
```

```
## Confusion Matrix and Statistics
##
```

```

##
## test_pred_10  1  2  3  4  5
##              1 18  3  1  2  0
##              2  3 18  0  7  0
##              3  2  0 20  2  1
##              4  3  3  5 14  0
##              5  0  1  0  0 24
##
## Overall Statistics
##
##              Accuracy : 0.7401575
##              95% CI : (0.6548765, 0.813901)
##              No Information Rate : 0.2047244
##              P-Value [Acc > NIR] : < 0.00000000000000022204
##
##              Kappa : 0.6752422
##              McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: 1  Class: 2  Class: 3  Class: 4  Class: 5
## Sensitivity          0.6923077 0.7200000 0.7692308 0.5600000 0.9600000
## Specificity          0.9405941 0.9019608 0.9504950 0.8921569 0.9901961
## Pos Pred Value       0.7500000 0.6428571 0.8000000 0.5600000 0.9600000
## Neg Pred Value       0.9223301 0.9292929 0.9411765 0.8921569 0.9901961
## Prevalence           0.2047244 0.1968504 0.2047244 0.1968504 0.1968504
## Detection Rate       0.1417323 0.1417323 0.1574803 0.1102362 0.1889764
## Detection Prevalence 0.1889764 0.2204724 0.1968504 0.1968504 0.1968504
## Balanced Accuracy    0.8164509 0.8109804 0.8598629 0.7260784 0.9750980
taulahog50<-(table(test_pred_50, max.col(dt_HOG_test[,names(dt_HOG_test)%in%resp_var_HOG])));
cmhog50<-confusionMatrix(taulahog50);cmhog50

## Confusion Matrix and Statistics
##
##
## test_pred_50  1  2  3  4  5
##              1 19  2  1  1  0
##              2  2 19  0  8  0
##              3  2  2 19  2  1
##              4  3  1  6 14  0
##              5  0  1  0  0 24
##
## Overall Statistics
##
##              Accuracy : 0.7480315
##              95% CI : (0.6633077, 0.8208324)
##              No Information Rate : 0.2047244
##              P-Value [Acc > NIR] : < 0.00000000000000022204
##
##              Kappa : 0.6850833
##              McNemar's Test P-Value : NA
##
## Statistics by Class:
##

```

```
##           Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.7307692 0.7600000 0.7307692 0.5600000 0.9600000
## Specificity      0.9603960 0.9019608 0.9306931 0.9019608 0.9901961
## Pos Pred Value   0.8260870 0.6551724 0.7307692 0.5833333 0.9600000
## Neg Pred Value   0.9326923 0.9387755 0.9306931 0.8932039 0.9901961
## Prevalence       0.2047244 0.1968504 0.2047244 0.1968504 0.1968504
## Detection Rate   0.1496063 0.1496063 0.1496063 0.1102362 0.1889764
## Detection Prevalence 0.1811024 0.2283465 0.2047244 0.1889764 0.1968504
## Balanced Accuracy 0.8455826 0.8309804 0.8307312 0.7309804 0.9750980
```

```
taulaex3<-(table(test_pred_ex3, dt_ex3_test[, 'label']));
cmex3<-confusionMatrix(taulaex3); cmex3
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##
```

```
## test_pred_ex3 helicopter kangaroo menorah starfish trilobite
```

```
##   helicopter      13         3         3         7         1
##   kangaroo        5         11        9         5         0
##   menorah         2          5       11         1         0
##   starfish        6          5         3        12         0
##   trilobite       0          1         0         0        24
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.5590551
```

```
##           95% CI : (0.4682592, 0.6470337)
```

```
##           No Information Rate : 0.2047244
```

```
##           P-Value [Acc > NIR] : < 0.0000000000000022204
```

```
##
```

```
##           Kappa : 0.4490239
```

```
##   McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##           Class: helicopter Class: kangaroo Class: menorah
```

```
## Sensitivity      0.5000000      0.4400000      0.42307692
```

```
## Specificity      0.8613861      0.81372549      0.92079208
```

```
## Pos Pred Value   0.4814815      0.36666667      0.57894737
```

```
## Neg Pred Value   0.8700000      0.85567010      0.86111111
```

```
## Prevalence       0.2047244      0.19685039      0.20472441
```

```
## Detection Rate   0.1023622      0.08661417      0.08661417
```

```
## Detection Prevalence 0.2125984      0.23622047      0.14960630
```

```
## Balanced Accuracy 0.6806931      0.62686275      0.67193450
```

```
##
```

```
##           Class: starfish Class: trilobite
```

```
## Sensitivity      0.48000000      0.9600000
```

```
## Specificity      0.86274510      0.9901961
```

```
## Pos Pred Value   0.46153846      0.9600000
```

```
## Neg Pred Value   0.87128713      0.9901961
```

```
## Prevalence       0.19685039      0.1968504
```


```
## Detection Rate   0.09448819      0.1889764
```

```
## Detection Prevalence 0.20472441      0.1968504
```

```
## Balanced Accuracy 0.67137255      0.9750980
```

```
data.frame(accuracy=c(cmhog10$overall[1],cmhog50$overall[1],cmex3$overall[1]),
  CPU_seconds=c(HOG_10_CPU_time,HOG_50_CPU_time, ex3_fit_time),
  model=c("HOG_10","HOG_50","NN_2_Hidden")) %>%
  melt(id.var="model") %>%
  ggplot(aes(x=variable,
    y=value,
    fill=model)) +
  geom_bar(stat='identity',
    position="dodge") +
  facet_wrap(~variable,scales="free")
```

```
## Warning: attributes are not identical across measure variables; they will
## be dropped
```



We believe that the three models are not comparable. The HOG ones are, but comparing them to the NN with 2 hidden layers makes less sense, This is due to: - The format of the images is not the same. HOG is a vector of 81 elements, whereas the neural network has a longer length: 1024. Therefore, HOG_50 takes longer than HOG_10 as it is more complex; but comparing it with the other alternative makes less sense. The Neural Network is the one that takes longer to compute. On the other hand, Neural Network with caret implements a 5 fold CV, whereas the raw implementation of neuralnet doesn't.

PERFORMANCE: In terms of performance, HOG_10 has the best one.

CNN architectures:

Categories: airplanes (800), Motorbikes (798) and Faces (435) are large categories. For each category reshape grayscale images to 32x32. Split the dataset in two halves (training and test). From the training part, fit a CNN according with the model

```
categories_ex5 <- c("airplanes","Motorbikes","Faces")

obtain_images_to_keras <- function(x,vector_categories){

  #gets the list of the files in a given directory first.
  path_imgs <- list.files(paste0(images_dir,vector_categories[x]),
                        pattern = ".jpg",
                        full.names = TRUE)
  #creates an auxiliary matrix where we will place the results.
  #every row is a image. it first contains the label, and then the 1024 vector
  dt_img<-list()
  #we iterate on all the images
  for(i in 1:length(path_imgs)) {
    img <- image_to_grey(path_imgs[i]) %>%
      img2arr()
    dt_img [[i]]<- img
  }
  dt_img <- array(unlist(dt_img), dim = c(length(path_imgs), 32, 32))
  return(dt_img)
}

airplanes.array<-obtain_images_to_keras(1,categories_ex5)
dim(airplanes.array)
motorbikes.array<-obtain_images_to_keras(2,categories_ex5)
dim(motorbikes.array)
faces.array<-obtain_images_to_keras(3,categories_ex5)
dim(faces.array)

dt_ex5 <- abind(obtain_images_to_keras(1,categories_ex5),
               obtain_images_to_keras(2,categories_ex5),along=1)
dt_ex5 <- abind(dt_ex5,obtain_images_to_keras(3,categories_ex5),along=1)

labels_ex5 <- c(rep(categories_ex5[1],number_images[categories==categories_ex5[1]]),
               rep(categories_ex5[2],number_images[categories==categories_ex5[2]]),
               rep(categories_ex5[3],number_images[categories==categories_ex5[3]])) %>%
  as.factor() %>%
  as.numeric

index <- sample(1:dim(dt_ex5)[1],(2/3)*dim(dt_ex5)[1])

x_train <- dt_ex5[index,,]
x_test  <- dt_ex5[-index,,]

x_train <- array_reshape(x_train, c(nrow(x_train), 32, 32 ,1))
x_test  <- array_reshape(x_test,  c(nrow(x_test),  32, 32 ,1))
```

```

y_train<- array(labels_ex5[index])
y_test<- array(labels_ex5[-index])
y_train <- to_categorical(y_train-1, 3)
y_test <- to_categorical(y_test-1, 3)

```

```

input_shape<-c(32, 32, 1)

```

```

#checking the dimensions
dim(x_train)

```

5. Convolution layer with: . filters = 20, kernel_size = c(3,3), activation = 'relu', . Pooling layer with: pool_size = c(2, 2) . Fully connected layer with: units = 128, activation = 'relu' . Output layer with: units = num_classes, activation = 'softmax' . Assess the performance of the CNN predicting the categories of test images and obtain the confusion matrix.

```

num_classes<-3
start5time<-Sys.time()
##Model Convolutional Neural Network. documentation from --> https://keras.io/
###En primer lloc cal instal·lar Anaconda. En segon lloc executes: library(keras) , install_keras(). Un c
input_shape<-c(32,32,1)
time<-Sys.time()
model5 <-keras_model_sequential()%>%
layer_conv_2d(filters=20, kernel_size=c(3,3),activation='relu', input_shape = input_shape)%>%
layer_max_pooling_2d(pool_size=c(2,2))%>%
layer_flatten()%>%
layer_dense(units=128, activation='relu')%>%
layer_dense(units=num_classes, activation='softmax')

#opt<-optimizer_adam( lr= 0.0001 , decay = 1e-6 )
model5 %>% compile(
  loss = 'categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = 'accuracy'
)

summary(model5)

model5cnnnet <- model5 %>% fit(
  x_train, y_train,
  epochs = 20, batch_size = 128,
  validation_split = 0.2
)
end5time<-Sys.time()
plotmodel5cnnnet<-plot(model5cnnnet);plotmodel5cnnnet

model5cnnnettest <- predict(model5, x_test)

```

6. From the training part, fit a CNN according with a model proposed by yourself. . Describe the model. . Assess the performance of the CNN predicting the categories of test images and obtain the confusion matrix . filters = 30, kernel_size = c(3,3), activation = 'tanh', . Pooling layer with: pool_size = c(3, 3) . Fully connected layer with: units = 128, activation = 'tanh' . Output layer with: units = num_classes, activation = 'sigmoid' . Assess the performance of the CNN predicting the categories of test images and obtain the confusion matrix. **

```

start6time<-Sys.time()
input.shape<-c(32,32,1)
time<-Sys.time()
model6 <-keras_model_sequential()%>%
layer_conv_2d(filters=30, kernel_size=c(3,3),activation='tanh', input_shape = input.shape)%>%
layer_max_pooling_2d(pool_size=c(3,3))%>%
layer_flatten()%>%
layer_dense(units=128, activation='tanh')%>%
layer_dense(units=3, activation='sigmoid')

#opt<-optimizer_adam( lr= 0.0001 , decay = 1e-6 )
model6 %>% compile(
  loss = 'categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = 'accuracy'
)

summary(model6)

model6cnnet <- model6 %>% fit(
  x_train, y_train,
  epochs = 12, batch_size = 128,
  validation_split = 0.2
)
end6time<-Sys.time()
plotmodel6cnnet<-plot(model6cnnet);plotmodel6cnnet

model6cnnettest <- predict(model6, x_test)

```

7. Compare these two CNN model at least in two characteristics: the performance and CPU time.

```

timemodel5cnnet<-end5time - start5time ;timemodel5cnnet
timemodel6cnnet<-end6time - start6time ;timemodel6cnnet

loss5<-summary(model5cnnet$metrics$loss);loss5
loss6<-summary(model6cnnet$metrics$loss);loss6

accuracy5<-summary(model5cnnet$metrics$acc);accuracy5
accuracy6<-summary(model6cnnet$metrics$acc);accuracy6

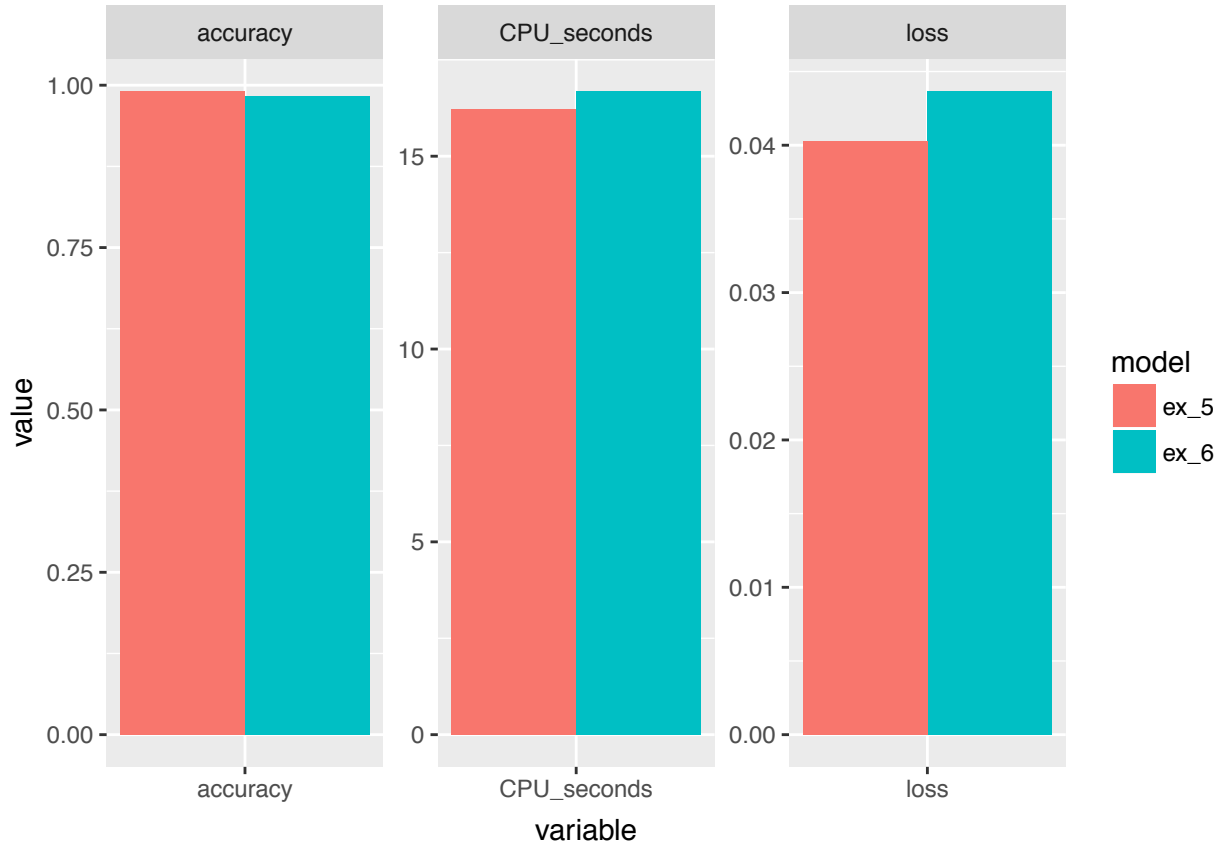
imgs_to_load<-list.files("objects_to_load/objectesR")
for (i in 1:length(imgs_to_load)){
  load(paste0("objects_to_load/objectesR/",imgs_to_load[i]))
}

data.frame(accuracy=c(accuracy5,accuracy6),
            CPU_seconds=c(timemodel5cnnet,timemodel6cnnet),
            loss=c(loss5,loss6),
            model=c("ex_5","ex_6")) %>%
melt(id.var="model") %>%
ggplot(aes(x=variable,
            y=value,
            fill=model)) +
geom_bar(stat='identity',

```

```
position="dodge") +
facet_wrap(~variable,scales="free")
```

```
## Warning: attributes are not identical across measure variables; they will
## be dropped
```



The best CNN model is the model5cnnnet. Although its CPU time is slightly higher than in model6cnnnet, its performance is better. It also has a lower loss.

8. Compare the NN (in 3) with the CNN (in 5) in the analysis of airplanes, Motorbikes and Faces categories.

```
categories_ex5 <- c("airplanes","Motorbikes","Faces")
dt_ex8 <- rbind(
  obtain_images(1, categories_ex5),
  obtain_images(2, categories_ex5),
  obtain_images(3, categories_ex5))

trainIndex <- createDataPartition(dt_ex8$label, p=.7, list=F)
dt_ex8_train <- dt_ex8[trainIndex, ]
dt_ex8_test <- dt_ex8[-trainIndex, ]

xnam <- names(dt_ex8_train[-1])
resp_var <- paste0(levels(dt_ex8_train$label), collapse=" + ")
(fmla <- as.formula(paste(resp_var, "~", paste(xnam, collapse=" +"))))
```

```

my.grid <- expand.grid(size=seq(from=1,to=5,by=3),
                      decay=seq(from=0.1,to=0.9,by=0.3))
ex_8_grid_search_time<-Sys.time()
ex8.fit <- train(label ~ .,
                 data = dt_ex8_train,
                 method = "nnet",
                 tuneGrid = my.grid,
                 MaxNWts=6036,
                 trControl = trainControl(method = "cv",
                                           verboseIter = FALSE,
                                           number = 2,
                                           returnData = FALSE),
                 verbose=FALSE)
ex_8_grid_search_time<-Sys.time() - ex_8_grid_search_time

save(ex_8_grid_search_time, file="objects_to_load/objectesR/ex_8_grid_search_time.RData")

#for the optimized parameters
my.grid <- expand.grid(size=ex8.fit$bestTune$size,
                      decay=ex8.fit$bestTune$decay)
ex_8_fit_time <- Sys.time()
ex8.fit <- train(label ~ .,
                 data = dt_ex8_train,
                 method = "nnet",
                 tuneGrid = my.grid,
                 MaxNWts=6036,
                 trControl = trainControl(method = "cv",
                                           verboseIter = FALSE,
                                           number = 2,
                                           returnData = FALSE),
                 verbose=FALSE)
ex_8_fit_time <- Sys.time() - ex_8_fit_time

# Plot model and predict
#plotnet(ex8.fit)
train_pred_ex8 <- predict(ex8.fit, newdata=dt_ex8_train[,-1], type="raw")
test_pred_ex8 <- predict(ex8.fit, newdata=dt_ex8_test[,-1], type="raw")

imgs_to_load<-list.files("objects_to_load/objectesR")
for (i in 1:length(imgs_to_load)){
  load(paste0("objects_to_load/objectesR/",imgs_to_load[i]))
}

taulaex8<-(table(test_pred_ex8, dt_ex8_test[,1]));
cmex8<-confusionMatrix(taulaex8); sum(diag(cmex8$table))/sum(cmex8$table)

save(ex_8_fit_time, file="objects_to_load/objectesR/ex_8_fit_time.RData")
save(cmex8, file="objects_to_load/objectesR/cmex8.RData")

```

Doing the grid search is extremely costly for the NN in the nnet integration of caret. It took 4.1993806163 to do the grid.search, and after that 1.5317859491 to just fit the model for the optimized parameters. It was also limited using the parameter MaxNWts, otherwise the computation was taking really long.

```

imgs_to_load<-list.files("objects_to_load/objectesR")
for (i in 1:length(imgs_to_load)){
  load(paste0("objects_to_load/objectesR/",imgs_to_load[i]))
}

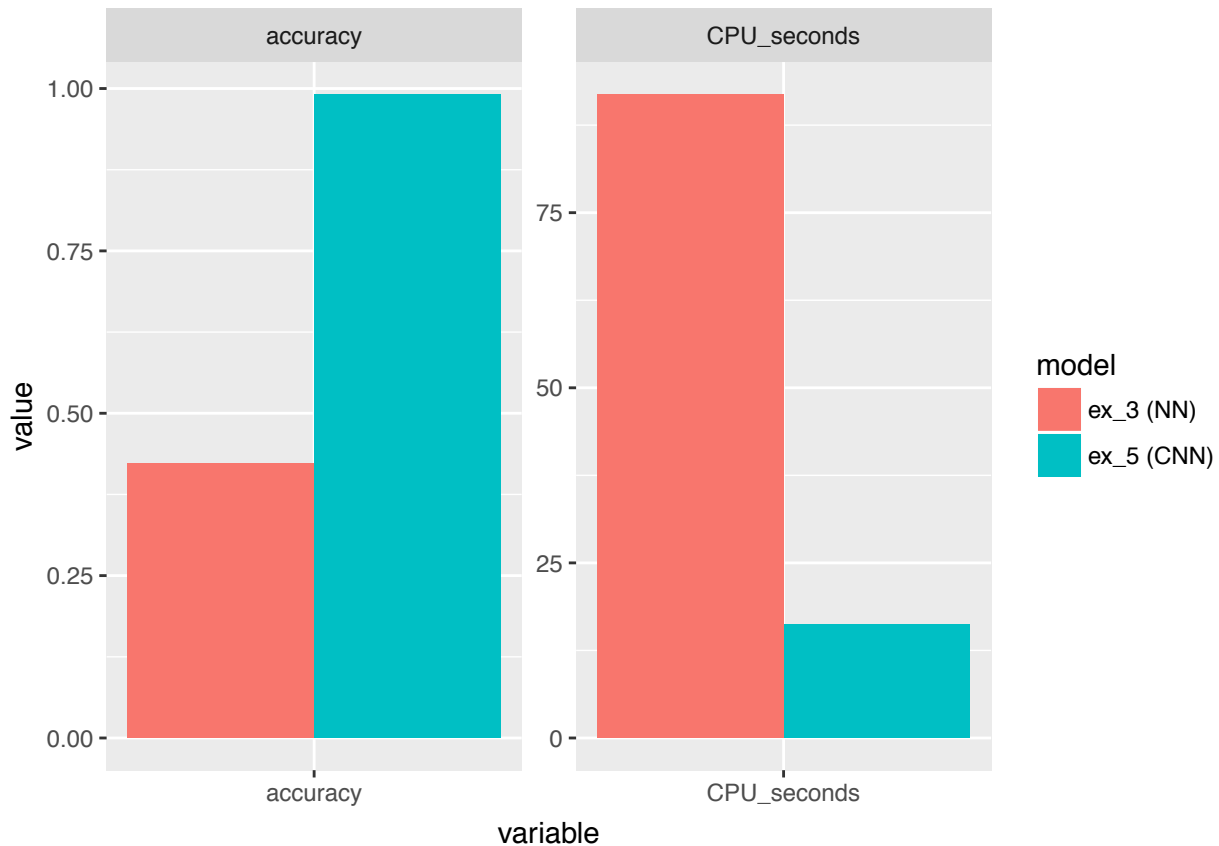
data.frame(accuracy=c(accuracy5,sum(diag(cmex8$table))/sum(cmex8$table)),
             CPU_seconds=c(timemodel5cnnnet,ex_8_fit_time),
             model=c("ex_5 (CNN)","ex_3 (NN)") %>%
melt(id.var="model") %>%
ggplot(aes(x=variable,
           y=value,
           fill=model)) +
geom_bar(stat='identity',
         position="dodge") +
facet_wrap(~variable,scales="free")

```

```

## Warning: attributes are not identical across measure variables; they will
## be dropped

```



Both performance wise (accuracy) and in terms of CPU time, the CNN is way better than the NN.